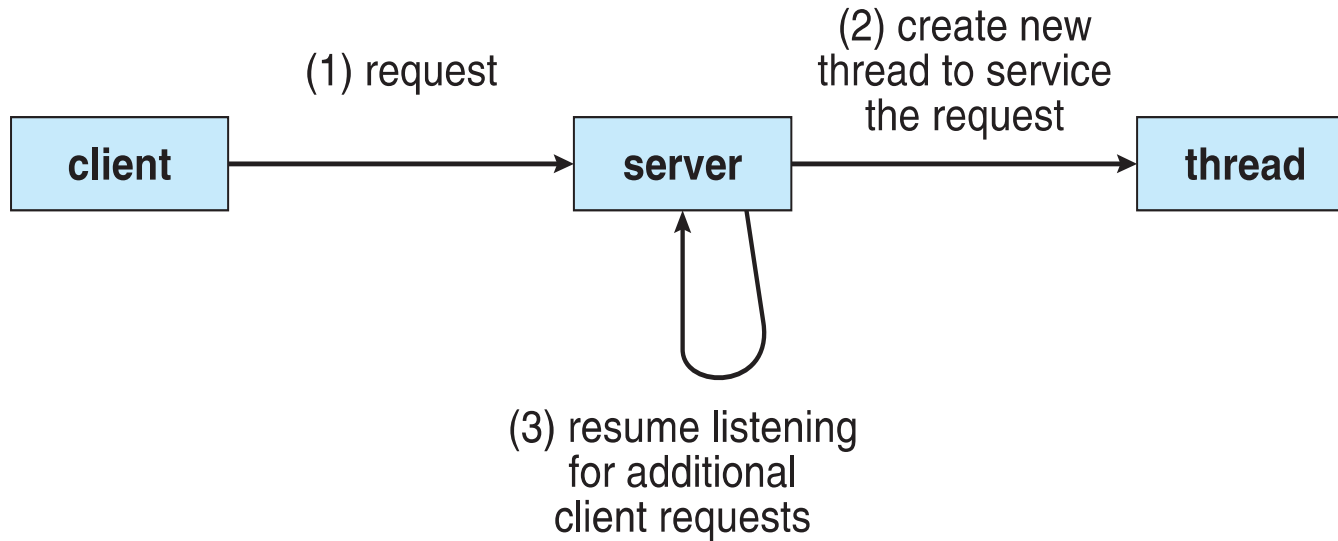


## **Chapter 4 -- Threads**

**A/Prof Jian Zhang**  
**School of Computing and Communication**  
**University of Technology, Sydney**

Ref: A. Silberschatz, P. B. Galvin & G. Gagne

# Multithreaded Server Architecture



# Benefits

- **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- **Scalability** – process can take advantage of multiprocessor architectures

\* **Scalability:** In a multiprocessor architecture, where threads may be running in parallel on different processing cores. A single-threaded process can run on only one processor, regardless how many are available. In such scenarios where a program might suffer from frequent page faults or has to wait for other system events, a multithreaded solution would perform better even on a single-processor system .

\* A **page fault** occurs when a process accesses a **page** that is mapped in the virtual address space, but not loaded in physical memory

# Multicore Programming


- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:
  - **Dividing activities**
  - **Balance**
  - **Data splitting**
  - **Data dependency**
  - **Testing and debugging**

- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
  - Single processor / core, scheduler providing concurrency
- .....


Notice: the distinction between **parallelism** and **concurrency** in this discussion. A system is parallel if it can perform more than one task **simultaneously**. In contrast, a concurrent system supports more than one task by allowing all the tasks to make progress. Thus, it is possible to have concurrency without parallelism.

# Concurrency vs. Parallelism

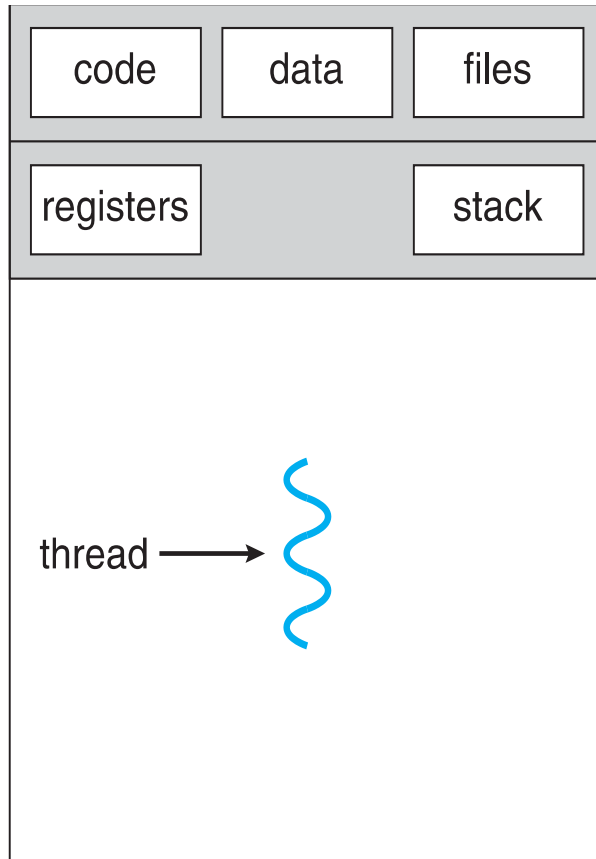
## ■ Concurrent execution on single-core system:



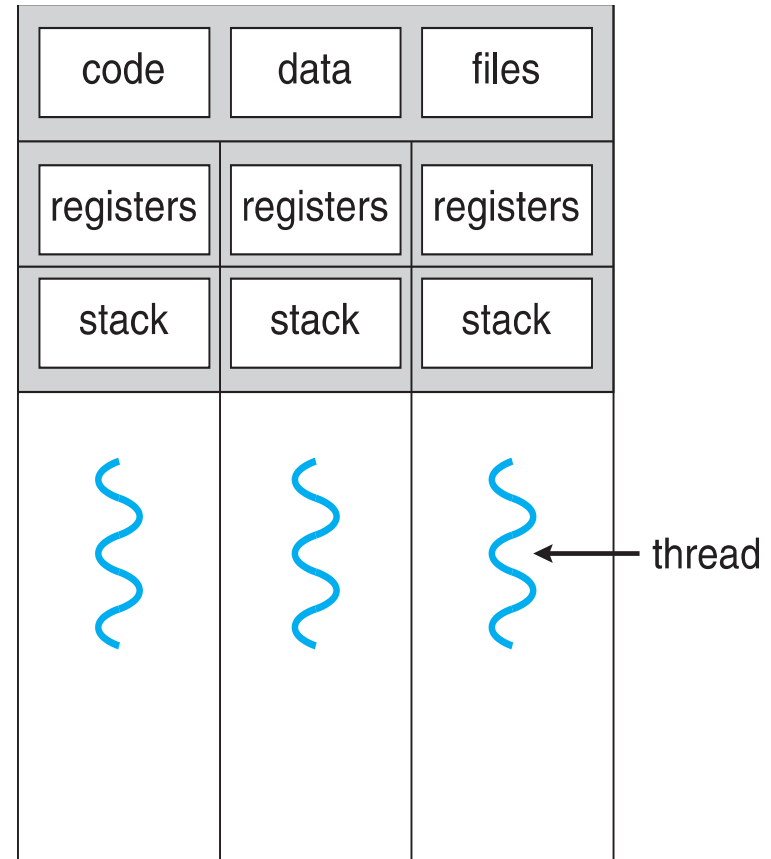
## ■ Parallelism on a multi-core system:



# Single and Multithreaded Processes



single-threaded process



multithreaded process

# Single and Multithreaded Processes


- A single-threaded: One by one process in programming.

Concept: Sequential program such as the next step relies on the previous step closely.

Example : Word process is better to have one thread to have editing, saving sequentially. Another example is a video paly program frame-by-frame!

# Many-to-One (Multithreading Models)

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multi-core system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
  - **Solaris Green Threads**
  - **GNU Portable Threads**






# One-to-One (Multithreading Models)

- Each user-level thread maps to kernel thread
- Creating a user-level thread and also creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead


- Examples

- Windows
- Linux
- Solaris 9 and later



# Many-to-Many Model (Multithreading Models)

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows with the *ThreadFiber* package



# Consider the following code segment:

## Check point

```
pid_t pid;
pid = fork();
if (pid == 0) { /* child process */
    //fork();
    thread create( . . . );
}
fork();
```

- a. How many unique processes are created?
- b. How many unique threads are created?

# Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
  - Library entirely in user space
  - Kernel-level library supported by the OS
- **Pthreads** may be provided either as user- or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ***Specification***, not ***implementation***
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)

# Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) { // Argument number should be 2
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) { // convert string to integer
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

# Pthreads Example (Cont.)

```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);

printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

# Program state shared across threads in a multithreaded process

## Program state -- Global variables

```
int sum; /* this data is shared by the thread(s) */  
void *runner(void *param); /* threads call this function */
```

## Program state – Heap memory

Each thread gets a stack, while there's typically only one heap for the application

Conclusion: each thread has its separate set of register values and a separate stack memory.

# Signal Handling

- n **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- n A **signal handler** is used to process signals
  1. Signal is generated by particular event
  2. Signal is delivered to a process
  3. Signal is handled by one of two signal handlers:
    1. default
    2. user-defined
- n Every signal has **default handler** that kernel runs when handling signal
  - | **User-defined signal handler** can override default
  - | For single-threaded, signal delivered to process



# Thread Cancellation

- Terminating a thread before it has finished
- Thread to be canceled is **target thread**
- Two general approaches:
  - **Asynchronous cancellation** terminates the target thread immediately
  - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);
```

# Linux Threads

- Linux (OS, system call) refers to them as ***tasks*** rather than ***threads***
- Thread creation is done through **clone()** system call
- **clone()** allows a child task to share the address space of the parent task (process)
  - Flags control behavior
- **struct task\_struct** points to process data structures (shared or unique)

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

# **Chapter 5 – Process Synchronization**

**A/Prof Jian Zhang**  
**School of Computing and Communication**  
**University of Technology, Sydney**

Ref: A. Silberschatz, P. B. Galvin & G. Gagne

# Background

- Processes can execute concurrently
  - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Illustration of the problem:

Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

# Race Condition

- **counter++ in Producer** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter-- in Consumer** could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute	<code>register1 = counter</code>	{register1 = 5}
S1: producer execute	<code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute	<code>register2 = counter</code>	{register2 = 5}
S3: consumer execute	<code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute	<code>counter = register1</code>	{counter = 6}
S5: consumer execute	<code>counter = register2</code>	{counter = 4}

# Critical Section Problem

- Consider system of  $n$  processes  $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code
  - Process may be changing common variables, updating table, writing file, etc
  - When one process in critical section, no other may be in its critical section
- ***Critical section problem*** is to design protocol to solve this
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

# Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non- preemptive

- **Preemptive** – allows preemption of process when running in kernel mode
- **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
  - Essentially free of race conditions in kernel mode

# Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.
- All solutions below based on idea of **locking**
  - Protecting critical regions via locks
- A single-processor system – to disable interrupts!
  - It can disable the timer interrupt and prevent context switching from taking place
  - Currently running code would execute without preemption
  - Generally, **too inefficient on multiprocessor systems**
    - Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
  - **Atomic** = non-interruptible
    - Either test memory word and set value
      - Function call -- test\_and\_set(&lock)
    - Or swap contents of two memory words
      - Function call -- compare and swap(int \*value, int expected, int new value)



# Mutex Locks

Previous solutions are complicated and generally inaccessible to application programmers

OS designers build software tools to solve critical section problem

Simplest is mutex lock

Protect a critical section by first **acquire()** a lock then **release()** the lock

Boolean variable indicating if lock is available or not

Calls to **acquire()** and **release()** must be atomic

Usually implemented via hardware atomic instructions

But this solution requires busy waiting

This lock therefore called a **spinlock\***

\*spin in place" waiting on the lock to be released by the first process, thus the name spin lock.

# Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to **synchronize their activities**.
- Semaphore **S** – integer variable \*
- Can only be accessed via two **indivisible (atomic) operations**
  - **wait()** and **signal()**
    - Originally called **P()** and **V()**
- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait!  
    S--;  
}
```
- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```

\* As usual synchronization primitive, a semaphore is based on a variable. This variable may be incremented or decremented and its state will represent ability to acquire lock.

# Example discussion

## Check point

- Semaphores as a server to limit the number of concurrent connections
  - Set up the number of allowable open socket connections
  - Acquire method is to set up a connection, Release method is to release the connection
  - When the connection reaches the limited number (set by semaphores), semaphore will block the acquire method until a new connection (released) is available

**5.23** Show how to implement the wait() and signal() semaphore operations in multiprocessor environments using the test and set() instruction. The solution should exhibit minimal **busy waiting**.

```
int guard = 0;
int semaphore value = 0;
wait()
{
    while (TestAndSet(&guard) == 1);
    if (semaphore value == 0) {
        atomically add process to a queue of processes
        waiting for the semaphore and set guard to 0;
    }else {
        semaphore value--;
        guard = 0;
    }
}

signal()
{
    while (TestAndSet(&guard) == 1);
    if (semaphore value == 0 && there is a process on the wait queue)
        //wake up the first process in the queue
        //of waiting processes (busy waiting)
    else
        semaphore value++;
    guard = 0;
}
```

# Busy waiting -- Check point

**busy waiting:** While a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the call to `acquire()`. It is waiting for a condition to be satisfied in a tight loop without relinquishing the processor

Busy waiting wastes CPU cycles that some other process might be able to use productively.

Alternatively, a process could wait by relinquishing the processor, and block on a condition and wait to be awakened at some appropriate time in the future

Busy waiting can be avoided but incurs the overhead associated with putting a process to sleep and having to wake it up when the appropriate program state is reached.

# Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
  - Same as a **mutex lock**
- Can solve various synchronization problems
- Consider  $P_1$  and  $P_2$  that require  $S_1$  to happen before  $S_2$   
Create a semaphore “**synch**” initialized to 0

**P1:**

$S_1;$

**signal(synch);**

**P2:**

**wait(synch);**

$S_2;$

- Can implement a counting semaphore  $S$  as a binary semaphore

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$   
`wait(S) ;`  
`wait(Q) ;`  
`...`  
`signal(S) ;`  
`signal(Q) ;`

$P_1$   
`wait(Q) ;`  
`wait(S) ;`  
`...`  
`signal(Q) ;`  
`signal(S) ;`

- **Starvation – indefinite blocking**
  - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock while needed by higher-priority process. it is possible that a third task  $M$  of medium priority will jump in ....
  - Solved via **priority-inheritance protocol**

# Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem

## *The bounded-buffer problem was introduced in Section 5.1*

- The producer and consumer processes share the following data structures:
  - Int ***n*** buffers, each can hold one item
  - Semaphore ***mutex*** initialized to the value 1
  - Semaphore ***full*** initialized to the value 0
  - Semaphore ***empty*** initialized to the value ***n***

Examples are in Figures 5.9 (the producer) and 5.10 (the consumer)

### *The bounded-buffer:*

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
    item; item buffer[BUFFER_SIZE];
    int in = 0; int out = 0;
```

\* Solution is correct, but can only use BUFFER\_SIZE-1 elements

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    //if (((in + 1) % BUFFER_SIZE) == out)
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter ++;
}
item next_consumed;
while (true) {
    //if (in == out)
    While (counter == 0)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter --;
    /* consume the item in next consumed */
}
```

**\*They may not function correctly when executed concurrently – counter == ???**



# Classical Problems of Synchronization

- `do {mutex=1; empty =1; full=0`
- `...`
- `/* produce an item in next produced */`
- `...`
- `wait(empty); // value n - 1`
- `wait(mutex); // =1 - 1`
- `...`
- `/* add next produced to the buffer */`
- `...`
- `signal(mutex); //release Mutex`
- `signal(full); // full: = + 1`
- `} while (true);`
- **Figure 5.9** The structure of the producer process.

- `do {mutex=1; empty =1; full=0`
- `....`
- `wait(full); // full= -1`
- `wait(mutex); //= 1-1`
- `...`
- `/* remove an item from buffer to next consumed */`
- `...`
- `signal(mutex); // release Mutex`
- `signal(empty); //value ++1`
- `...`
- `/* consume the item in next consumed */`
- `...`
- `} while (true);`
- **Figure 5.10** The structure of the consumer process

# Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write
- Problem – allow multiple readers to read at the same time
  - Only one single writer can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
  - Data set
  - Semaphore **rw\_mutex** initialized to 1
  - Semaphore **mutex** initialized to 1
  - Integer **read\_count** initialized to 0

# Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- **Abstract data type -- ADT**, internal variables only accessible by code within the procedure
- ADT -- encapsulates data with a set of functions to operate on that data that are independent of any specific implementation of the ADT.
- Only one process may be active within the monitor at a time
- But not powerful enough to model some synchronization schemes

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }

    procedure Pn (...) {.....}

    Initialization code (...) { ... }
}
}
```

# Condition Variables

An additional synchronization mechanisms –  
Condition construct

## **Condition $x, y$ ;**

- Two operations are allowed on a condition variable:
  - **$x.\text{wait}()$**  – a process that invokes the operation is suspended until  **$x.\text{signal}()$**
  - **$x.\text{signal}()$**  – resumes exactly one of suspended processes (if any) that invoked  **$x.\text{wait}()$**

# Condition Variables Choices

- If process P invokes **`x.signal()`** , and process Q is suspended in **`x.wait()`** , what should happen next?
  - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Otherwise, both P and Q would be active simultaneously within the monitor. The options include
  - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
  - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
  - Both have pros and cons – language implementer can decide
  - Monitors implemented in Concurrent Pascal compromise
    - P executing signal immediately leaves the monitor, Q is resumed
  - Implemented in other languages including Mesa, C#, Java

# Linux Synchronization

- Linux:
  - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
  - Version 2.6 and later, fully preemptive
- Linux provides:
  - Semaphores
  - atomic integers
  - spinlocks
  - reader-writer versions of both
- On single-cpu system, spinlocks replaced by enabling and disabling kernel preemption

## **/\*Ruuner + Mutex + Semaphore\*/**

```
#include <stdlib.h>
#include <stdio.h>
#include <pthread.h>
#include <semaphore.h>

int sum; /*this data is shared by the threads*/
/* The mutex lock */
pthread_mutex_t mutex;
/* the semaphores */
sem_t one, two;
pthread_t tid1, tid2;          //Thread ID
pthread_attr_t attr; //Set of thread attributes

void *runnerOne(void *param); /*thread one call this function*/
void *runnerTwo(void *param); /*thread two call this function*/
void initializeData();

int main(int argc, char*argv[])
{
    // pthread_t tid; /*the thread identifier*/
    // pthread_attr_t attr; /*set of thread attributes*/

    if(argc!=2){
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if(atoi(argv[1])<0){
        fprintf(stderr, "%d must be >=0\n", atoi(argv[1]));
        return -1;
    }
}
```

```

initializeData();

/*get the default attributes*/
pthread_attr_init(&attr);

pthread_create(&tid1,&attr,runnerOne,argv[1]); /*create the thread*/
printf("sum=%d\n",sum);

pthread_create(&tid2,&attr,runnerTwo,argv[1]); /*create the thread*/

sem_post(&two); /* unlock the 'two' semaphore */

//sem_post(&one); /* unlock the 'one' semaphore */

        /*wait for the thread to exit*/
        pthread_join(tid1,NULL);
        pthread_join(tid2,NULL);
        printf("sum=%d\n",sum);
    }/* end of Main */

void initializeData() {
    /* Create the mutex lock */
    pthread_mutex_init(&mutex, NULL);

    /* Create the 'one' semaphore and initialize the value = 0 the subroutine
    call is blocked*/    sem_init(&one, 0, 0); //

    /* Create the 'two' semaphore and initialize the value = 0 the subroutine
    call is blocked */
    sem_init(&two, 0, 0);

    /* Get the default attributes */
    pthread_attr_init(&attr);
}

```



```
/*The thread one will begin control  
in this function*/
```

```
void *runnerOne(void *param)  
{  
    /* acquire the 'one' semaphore */  
    sem_wait(&one);  
    /* acquire the mutex lock */  
    pthread_mutex_lock(&mutex);  
  
    int i, upper=atoi(param);  
    sum=0;  
  
    for(i=0; i<=upper; i++)  
        sum+=2*i;  
  
    printf("thread one\n");  
    /* release the mutex lock */  
    pthread_mutex_unlock(&mutex);  
    /* release and unlock the 'two'  
semaphore */  
    sem_post(&two);  
  
}
```

```
/*The thread two will begin control  
in this function*/
```

```
void *runnerTwo(void *param)  
{  
    /* acquire the 'two' semaphore */  
    sem_wait(&two);  
    /* acquire the mutex lock */  
    pthread_mutex_lock(&mutex);  
  
    int i, upper=atoi(param);  
    sum=10;  
  
    for(i=0; i<=upper; i++)  
        sum+=i;  
    printf("thread two\n");  
    /* release the mutex lock */  
    pthread_mutex_unlock(&mutex);  
    /* release and unlock the 'one'  
semaphore */  
    sem_post(&one);  
}
```